

Notes and exercises from “An Introduction to Statistical Learning” by James, Witten, Hastie & Tibshirani

Nicholas McGlinchey

July 21st 2016

Contents

Chapter 1	1
-----------	---

Chapter 1

The introduction outlines three basic tasks in statistical learning:

1. **Regression** - predicting *continuous numerical* values from other continuous values
2. **Classification** - predicting class membership; has a *categorical* or *qualitative* output.

In both these tasks, we are trying to predict an output.

3. **Clustering** - here we are not trying to predict an output variable, but instead are trying to understand how similar/different a set of observations are, usually based on a number of variables.

It is clear that I need to revise matrix multiplication, I didn't really understand the example on p. 12.

The book's website is [here](#).